

# Real-Time Block Rate Targeting

Thomas M. Harding

tom@chain2.org

October 11, 2019

## Abstract

A proof-of-work blockchain uses a retargeting algorithm, also termed a difficulty adjustment algorithm, to manage the rate of block production in the presence of changing hashrate. To derive the parameters that guide the search for the next block, nearly all such algorithms rely on averages of past inter-block time observations, as measured by on-chain timestamps. We are motivated to seek better responsiveness to changing hashrate, while improving stability of the block production rate and retaining the progress-free property of the original Bitcoin retargeting algorithm. We describe a class of retargeting algorithms for which the sole inter-block time input is that of the block being searched for, and whose response is nonlinear in that time. We discuss how these algorithms allow the other consensus rules that govern allowable timestamps to be tightened, which may improve the blockchain's effectiveness as a time-stamping machine.

## 1 Introduction

### 1.1 Traditional retargeting algorithm

A traditional retargeting algorithm sets a hash target  $G_n$  that defines a valid cryptographic hash of block  $n$ . The hash is required to be no larger than the target:

$$H(\text{block } n) \leq G_n \quad (1)$$

which gives rise to the block's proof-of-work [1].  $G_n$  is a function of data found in earlier blocks and is committed to as part of the header of block  $n$ . It is set so as to expect future inter-block times of a desired value  $T$ , given the hashrate estimate derived from past inter-block times. In the case of Bitcoin, it is recalculated every 2016 blocks, and unchanged otherwise. Omitting some details inessential to our examination:

$$G_n = \begin{cases} G_{n-1}^{\frac{s_{(n-1)} - s_{(n-2016)}}{2016T}}, & n \bmod 2016 = 0 \\ G_{n-1}, & \text{otherwise} \end{cases} \quad (2)$$

where  $s_i$  is the committed timestamp of block  $i$  and  $T$  ( $= 600$  seconds) is the desired inter-block time.  $G_n$  is known as soon as block  $n - 1$  is published.

Equation 2 makes use of the scaling property of the exponential distribution [3]. Post-retarget, assuming constant hashrate since block  $n-1$ , a valid block  $n$  will ideally be found after inter-block time  $t$  that is exponentially distributed with rate

$$\lambda = \frac{1}{T} \quad (3)$$

$$t \sim \text{Exp}(\lambda) \quad (4)$$

which describes a homogeneous Poisson process. Of course, hashrate will not be constant on the blockchain of interest and, as a random variable,  $t$  exhibits sample variance, so retargeting continues *ad infinitum*. Retargeting itself makes the global block production process not homogeneous Poisson.[4]

## 1.2 Unresponsiveness of the traditional algorithm

The observed inter-block time  $t$  has a standard deviation as large as its mean, which likely drove the design decision to aggregate 2016 consecutive samples before retargeting Bitcoin.

That decision reduced the coefficient of variation from 100% to 2.2%, but has resulted in a >5% overshoot versus target of the long term block production rate in the environment of ever-increasing hashrate.

More recently, Bitcoin has shared the pool of global SHA256 hashrate with other blockchains such as Bitcoin Cash, which have been assigned value in the eyes of the market. As miners are free to allocate their hashrate, a dynamic situation has emerged wherein miners follow diverse strategies to allocate their hashrate among the blockchains where it can be applied.

In this context, the Bitcoin retargeting algorithm has a weakness. It is slow to react to hashrate changes that occur on the timescale of changes in the price of the token (BTC, BCH, etc.). For example, a factor-of-2 decrease in the price of BTC versus other blockchains risks slowing block production by the same factor until the next retargeting, as hashrate departs.

We do not attempt to survey the altcoin landscape or its myriad retargeting algorithms, but the phenomenon of coin-hopping is nothing new to its inhabitants. Swings in the varied types of altcoin hashrate have been dramatic enough to end the existence of the altcoin as its miners suddenly depart and blocks are produced so slowly that users lose patience altogether [5, 8].

Because Bitcoin Cash has a more dynamic retargeting algorithm with a 144-block interval, it reacts more quickly to relative price changes, adjusts its own block production, and mitigates the profitability disparity, to the benefit of both blockchains. We see the advantage of responsiveness and are prompted to ask how we can increase it.

## 1.3 The memoryless property vs. progress-freedom

The exponential is the unique *memoryless* statistical distribution [2, chapter 4]. Being memoryless is the property that the likelihood of observing any particular

wall-clock time interval from now until the appearance of the next block is independent of the time passed since the appearance of block  $n - 1$ , whether that time is zero, 48 seconds, or any other value.

*Progress-freedom* is distinct from being memoryless. We define progress-freedom to be the property that a miner’s chance of finding a block at any moment is independent of how much work he has already done since the appearance of block  $n - 1$ .

Progress-freedom is a critical property for a proof-of-work blockchain because it results in a miner’s chance to find the next block being proportional to his hashrate. If a miner were to somehow get credit for progressive work done toward a block solution, larger miners would increase their chances of finding the block with each passing second, at the expense of smaller miners, creating an incentive to centralize mining.

The exponential distribution is progress-free, but not uniquely so. We describe a system that remains progress-free, but aims to be more responsive to hashrate changes, to exhibit more stable inter-block times, and be more resistant to target and timestamp manipulation attacks, while necessarily discarding the memoryless property.

## 2 Engineering the block production rate

Consider a more general block production rate function

$$\lambda(t) = at^{k-1} \tag{5}$$

Equation 5 contemplates a system where, during the search for block  $n$ , the instantaneous block production rate is not constant, but rather a monomial function of the time since block  $n - 1$ . Equation 3 is a special case ( $a = 1/T$ ,  $k = 1$ ) of Equation 5.

For the simple reason that it depends only on time, and not on accumulated work, Equation 5 is progress-free. Time-dependent *rate functions* are well studied in statistics and are also known as *hazard rate functions* (the hazard is something like a machine failing).

We set  $k > 1$ , so that with each passing second, the instantaneous block production rate increases. However, this just as strongly affects the likelihood that the block has not already been found, its *survival rate*  $1 - F(t)$ . [Figure 1]

The net result is a quasi-symmetrical distribution of inter-block times [Figure 2], whose shape is narrower with higher values of  $k$ . We are free to choose  $k$ , but we also require the mean inter-block time to be  $T$ . Therefore, we are not free to choose  $a$ ; we derive it.

It turns out that the hazard rate function uniquely defines the distribution of observed times by the relation:

$$F(t) = 1 - \exp \left[ - \int_0^t \lambda(t) dt \right] \quad [2, \text{chapter 4.1}] \tag{6}$$




Figure 1: RTT density factors (log scale)

Substituting Equation 5 into Equation 6 and solving the integral gives the cumulative distribution function

$$F(t) = 1 - e^{-\frac{at^k}{k}} \quad (7)$$

and probability density function ( $= F'(t)$ )

$$f(t) = at^{k-1} e^{-\frac{at^k}{k}}$$

which is a Weibull distribution [2, chapter 5.2] (notably, Weibull reduces to exponential when  $k = 1$ ). The expected value, or mean, is

$$\begin{aligned} E[f(t)] &= \int_0^\infty at^k e^{-\frac{at^k}{k}} \\ &= \Gamma(1 + \frac{1}{k}) (\frac{k}{a})^{\frac{1}{k}} \end{aligned} \quad (8)$$

Setting Equation 8 equal to  $T$  and solving for  $a$  gives the value of  $a$  in terms of the constants  $k$  and  $T$ :

$$a = k \left[ \frac{\Gamma(1 + \frac{1}{k})}{T} \right]^k \quad (9)$$

### 3 Real-Time Targeting

To keep the blockchain verifiable,  $t$  must be the difference in committed timestamps between the block  $n$  being searched for and the previous block:

$$t = s_n - s_{(n-1)}$$

The hash target  $G_{(n-1)}$  of the previous block, starting with the genesis block, is known. This committed target continues to be interpreted according to Equation 1, namely as the maximum hash that would be valid *if* the search for the block  $n$  were carried out using the traditional algorithm.

But we introduce changes to the *actual* search algorithm, described in the next three subsections.

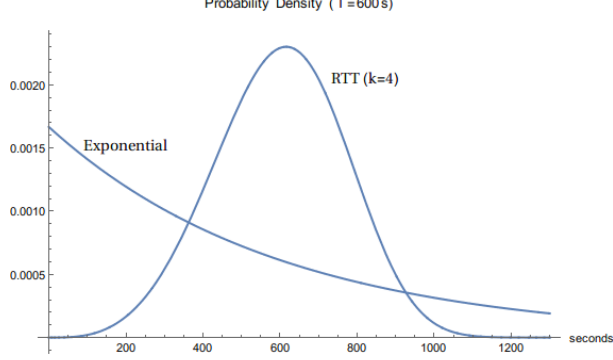


Figure 2: RTT density (compare exponential)

### 3.1 Subtargeting

To conduct the block search using the increasing block production rate of Equation 5, we define a *subtarget* at time  $t$ , such that

$$g_n(t) = G_{(n-1)} \frac{\lambda(t)}{\lambda} \quad (10)$$

whereby we have applied the desired block production rate, and "backed out" the traditional production rate. Combining Equations 3, 5, 9 and 10 gives

$$\begin{aligned} g_n(t) &= G_{(n-1)} k \frac{\Gamma(1 + \frac{1}{k})^k}{T^{k-1}} t^{k-1} \\ &= G_{(n-1)} \cdot \text{constant} \cdot t^{k-1} \end{aligned}$$

The proof-of-work validity constraint / mining target depends on  $t$  and is given by

$$H(\text{block } n) \leq g_n(t)$$

### 3.2 Modified retargeting

The subtargeting process results in a block found at time  $t$  which is distributed as shown in Figure 2. The quasi-symmetry and low (81% lower for  $k = 6$ ) variance compared to exponential are key. To a chosen degree,  $t$  is forced toward  $T$ , "calming down" the observed and recorded inter-block time and making it resilient to changed hashrate, before the next target is even computed.

However, the resilience of  $t$  works against us when retargeting for the next block. That operation must be exposed to the full effect of the hashrate implied by  $t$ . To achieve this is straightforward given the form of the function we chose for Equation 5. We find  $F(t)$ , also known as the *quantile* of  $t$ , and plug it into the inverse CDF of the exponential distribution, giving an equivalent time  $u$  in that space:

$$u = -\frac{\ln(1 - F(t))}{\lambda} \quad (11)$$

Combining Equations 3, 7, 9 and 11 gives

$$u = aT \frac{t^k}{k}$$

and each block header commits to a homogeneous-Poisson-equivalent target

$$G_n = G_{(n-1)} \frac{u}{T}$$

which is analogous to Equation 2, but which

- references the timestamp in block  $n$  itself:  $G_n \leftarrow u \leftarrow t \leftarrow s_n$ .
- uses only one inter-block time ( $t$ ) as input.
- is recomputed with every block.
- is not used to validate the hash of block  $n$  (see Section 3.1) (although the target itself is validated), but rather block  $n + 1$ .

Note that the subtarget  $g(t)$  does not influence  $G_n$  except through its role in the discovery of  $t$ .

Remarkably, there is no averaging whatsoever of past inter-block times. Instead, the subtargeting process is a kind of auction, which starts at a level difficult for the network to achieve, and continually becomes easier until the network achieves it with the hashrate available. For example, with  $k = 4$  and  $T = 600s$ , producing a block with  $t = 1s$  is some 80 million times more difficult than in the traditional system.

### 3.3 [Adjustment to $T$ ]

Until now, we have either chosen constants or derived them. No fitting or estimation has been done. By statistical simulation, and mining scenario simulation, we have verified that a random RTT-distributed variate does have mean  $T$  and the predicted variance, and statistics for the distribution of the transformed variable  $u$  are those of an exponential (homogeneous Poisson) with mean  $T$ , though we have not carried out these tests rigorously.

But we return to the statement made in Section 1.1, that even without our deliberately inhomogeneous construction, the act of periodic retargeting itself makes the global block production process not homogeneous.[4] This is true enough in the case of Bitcoin, with its 2016-block retargeting interval. But as we retarget every block, it cannot be ignored.

With constant hashrate, both mining scenario simulation and actual hashrate simulation show a mean inter-block time that is some 13% too long. This is correctable by introducing an exponential moving average of as few as 8 past block times, but we reject that course of action because it would compromise our primary goal of responsiveness to hashrate fluctuations.

Instead we simply adjust the desired mean inter-block time  $T$ , to a value of 531 seconds where unadjusted  $T = 600s$ . This adjustment corrects the constant-hashrate simulations to a mean very close to 600 seconds. Variable or trending

hashrate will produce a different mean, but this is true of every retargeting algorithm. The adjustment is only made to retargeting.

## 4 Parent chainwork

Since a block's chainwork now depends sensitively on its inter-block time  $t$ , the block comparator is changed to reference the chainwork of the parent, block  $n - 1$ , rather than the chainwork of the block itself.

This change ensures that blocks with the same parent have the same chainwork, and denies the opportunity for miners to conduct inexpensive "orphaning" attacks by starting a competing chain with a block whose timestamp is one second earlier than another miner's block.

## 5 Tighter timestamp validity rules

In Section 3.2, we mentioned that it is very difficult to mine a block with very small  $t$ . It is evident, too, that it is impossible to mine a block with  $t = 0$ , as Equation 5 makes plain that the block production rate is zero in that case.  $t < 0$  is just as nonsensical.

Consequently, it is natural that  $t > 0$  be enforced as a consensus rule. Furthermore, without negative inter-block timestamp differences, it may not be possible to recover from a situation where timestamps get ahead of the real-world clock. So it is very natural that we also tighten the maximum allowable timestamp to "now" instead of Bitcoin's "now + 2 hours". We doubt anyway that miners will be eager to accept a peer's "block from the future" that was easier to find than allowed right now.

How well the committed timestamps match actual time is an emergent result of the system, but since there is strong incentive to choose the latest possible non-future timestamp, i.e. a timestamp of "now", we have reason to be optimistic.

## 6 Conclusion

We have described a class of algorithms for retargeting the block production rate of a proof-of-work blockchain.

Our aim in the preceding sections has been to promote understanding of these algorithms when they are seen in operation. We have not attempted to formally specify a particular algorithm, to exhaustively quantify its responsiveness or overall performance, nor to compare its performance to other algorithms. These pursuits remain for the future.

## 7 Acknowledgements

We are particularly indebted to Scott Roberts [5], for the many insights and results stemming from his work across several altcoins; to jameslee777 [8, 9] for his work on Komodo, and to Andrew Stone [6, 7] for his seminal presentation of some of these ideas.

## References

- [1] Satoshi Nakamoto: Bitcoin: A peer-to-peer electronic cash system, 2008.
- [2] Sheldon Ross: A First Course in Probability, Macmillan Publishing Company, New York, 2nd edition, 1984.
- [3] Andre Nicolas: Is the family of exponential distributions closed under scaling?, Mathematics Stack Exchange, <https://math.stackexchange.com/q/85578>, version 2011-11-25.
- [4] R. Bowden, H.P. Keeler, A.E. Krzesinski and P.G. Taylor: Block arrivals in the Bitcoin blockchain, <https://arxiv.org/pdf/1801.07447>, 2018.
- [5] Scott Roberts: TSA: Change Difficulty During the Block, <https://github.com/zawy12/difficulty-algorithms/issues/36>, retrieved Sep. 2019.
- [6] Andrew Stone: The Blockchain Difficulty Information Paradox, <https://medium.com/@g.andrew.stone/the-blockchain-difficulty-information-paradox-879b0336864f>, retrieved Sep. 2019.
- [7] Andrew Stone: Tail Removal Block Validation, <https://medium.com/@g.andrew.stone/tail-removal-block-validation-ae26fb436524>, retrieved Sep. 2019.
- [8] jl777: AdaptivePoW: the solution to diff stranding of smaller blockchains, <https://medium.com/@jameslee777/adaptivepow-the-solution-to-diff-stranding-of-smaller-blockchains-425609df5563>, retrieved Sep. 2019.
- [9] jl777: Evolution of AdaptivePoW, <https://medium.com/@jameslee777/evolution-of-adaptivepow-dfea220d343f>, retrieved Sep. 2019.